

Data errors, how to find them?

Edwin de Jonge, Statistics Netherlands

@edwindjonge | github.com/edwindj



Who am I?


- ▶ Data scientist / Methodologist at Statistics Netherlands (aka CBS).
- ▶ Author of several R-packages, including `whisker`, `validate`, `errorlocate`, `docopt`, `daff`, `tableplot`, `ffbase`, `chunked`,
...
- ▶ Co-author of *Statistical Data Cleaning with applications in R (2018)* (together with @markvdloo)

Data cleaning...

A large part of your job is spent in data-cleaning:

- ▶ getting your data in the right shape (e.g. `tidyverse`, `dplyr`)
- ▶ assessing missing data (e.g. `VIM`, `datamaid`)
- ▶ checking validity (e.g. `validate`)
- ▶ locating and removing errors: **`errorlocate!`**
- ▶ impute values for missing or erroneous data (e.g. `simputation`, `VIM`, `recipes`)

Statistical Value Chain





**KEEP
CALM
AND
VALIDATE**

Validation rules?

Package validate allows to:

- ▶ formulate explicit data rule that data must conform to:

```
library(validate)
check_that( data.frame(age=160, driver_license=TRUE),
  age >= 0,
  age < 150,
  if (driver_license == TRUE) age >= 16
)
```

Explicit validation rules:

- ▶ Give a clear overview what the data must conform to.
- ▶ Can be used to reason about.
- ▶ Can be used to fix/correct data!
- ▶ Find error, and when found correct it.

Note:

- ▶ Manual fix is error prone, not reproducible and not feasible for large data sets.
- ▶ Large rule set have (very) complex behavior, e.g. entangled rules: adjusting one value may invalidate other rules.

Error localization

Error localization is a procedure that points out fields in a data set that can be altered or imputed in such a way that all validation rules can be satisfied.

Find the error:

```
library(validate)
check_that( data.frame(age=160, driver_license=TRUE),
  age >= 0,
  age < 150,
  if (driver_license == TRUE) age >= 16
)
```

It is clear that age has an erroneous value, but for more complex rule sets it is less clear.

Multivariate example:

```
check_that( data.frame( age      = 3
                        , married = TRUE
                        , attends = "kindergarten"
                        )
, if (married == TRUE) age >= 16
, if (attends == "kindergarten") age <= 6
)
```

Ok, clear that this is a faulty record, but what is the error?

Feligi Holt formalism:

Find the minimal (weighted) number of variables that cause the invalidation of the data rules.

Makes sense! (But there are exceptions...)

Implemented in `errorlocate` (second generation of `editrules`).

Formal description (1)

Rule $r_i(\mathbf{x})$

A rule a disjunction of atomic clauses:

$$r_i(\mathbf{x}) = \bigvee_j C_i^j(\mathbf{x})$$

with:

$$C_i^j(\mathbf{x}) = \begin{cases} \mathbf{a}^T \mathbf{x} \leq b \\ \mathbf{a}^T \mathbf{x} = b \\ x_j \in F_{ij} \text{ with } F_{ij} \subseteq D_j \\ x_j \notin F_{ij} \text{ with } F_{ij} \subseteq D_j \end{cases}$$

Rule system:

The rules form a system $R(\mathbf{x})$:

$$R_H(\mathbf{x}) = \bigwedge_i r_i$$

If $R_H(\mathbf{x})$ is true for record \mathbf{x} , then the record is valid, otherwise one (or more) of the rules is violated.

Mixed Integer Programming to FH

Each rule set $R(\mathbf{x})$ can be translated into a mip problem and solved.

$$\begin{aligned} &\text{Minimize } f(\mathbf{x}) = 0; \\ &\text{s.t. } \mathbf{R}\mathbf{x} \leq \mathbf{d} \end{aligned}$$

- ▶ $f(\mathbf{x})$ is the (weighted) number of changed variable: $\delta_i \in 0, 1$

$$f(\mathbf{x}) = \sum_{i=1}^N w_i \delta_i$$

- ▶ \mathbf{R} contains rules: $\mathbf{R}_H(\mathbf{x}) \leq \mathbf{d}_H$ and soft constraints: $\mathbf{R}_0(\mathbf{x}, \delta) \leq \mathbf{d}_0$ that try fix the values of \mathbf{x} to the measured values.

errorlocate

- ▶ translates your rules automatically into a mip form.
- ▶ Uses lpSolveAPI to solve the problem.
- ▶ contains a small framework for implementing your own error localization algorithms.

errorlocate::locate_errors

```
locate_errors( data.frame( age      = 3
                           , married = TRUE
                           , attends = "kindergarten"
                           )
              , validator( if (married == TRUE) age >= 16
                           , if (attends == "kindergarten") age <= 6
                           )
              )$errors
```

```
##           age married attends
## [1,] FALSE      TRUE  FALSE
```


errorlocate::replace_errors

```
replace_errors(  
  data.frame( age      = 3  
              , married = TRUE  
              , attends = "kindergarten"  
            )  
  , validator( if (married == TRUE) age >= 16  
              , if (attends == "kindergarten") age <= 6  
            )  
)
```

```
##   age married      attends  
## 1    3      NA kindergarten
```

Pipe %>% friendly


The `replace_errors` function is pipe friendly:

```
rules <- validator(age < 150)

data_noerrors <-
  data.frame(age=160, driver_license = TRUE) %>%
  replace_errors(rules)

errors_removed(data_noerrors) # contains errors removed
```

Interested?



SDCR

M. van der Loo and E. de Jonge
(2018) *Statistical Data Cleaning
with applications in R* Wiley, Inc.

errorlocate

► Available on [CRAN](#)

More theory?

← See book

Thank you for your attention (and enjoy The Hague)!